

Databases and Migrations

COSC 480

21 February 2018

Joel Sommers

`jsommers@colgate.edu`

Colgate University

The Database is Gold

- Contains valuable customer data — don't want to test your app on the *real* data!
- Rails solution: development, production, and test *environments* each have their own DB
 - Different DB types appropriate for each
 - E.g., simple file-based database for test/development, like sqlite3
 - E.g., Industry-strength DB for production, like mysql, postgresql, or something else
- How to make changes to each DB, and maintain consistency across environments?
- Rails solution: **migration**
 - **Migration**: a script describing changes to tables in the database

Migration Advantages

- Can identify each migration and know which ones have been applied, and when
 - Migrations can often be designed to be reversible
- Can manage with version control
- Automated == reliably repeatable!
 - Compare: use bundler vs. manually installing libraries/gems
- Theme: don't do it manually, automate it!
 - Specify what to do, create tools to automate

Meet a code generator

```
rails generate migration CreateProducts
```

- Create the migration template
- Note: it is only created, not *applied* yet
- Apply migration to development DB: `rails db:migrate`
- A table describing which migrations have been applied is part of the database!
- Apply migration to production DB:

```
heroku run:detached rails db:migrate
```

Applying migration also records in DB itself which migrations have been applied.

Adding a new or updating a) model to a Rails app

1. Create a migration describing the changes:

`rails generate migration (name)`

- Type `rails g migration` to get help
- Many naming conventions to make migrations easy

2. Apply the migration: `rails db:migrate`

3. If new model (instead of modifying an existing a model)

3.1 Create model file in `app/models/model.rb`

3.2 Update test DB schema: `rails db:test:prepare`

Can also use the rails model generator

- `rails generate model`
- It also creates a migration script along with `app/models/model.rb`
 - Can specify column names/types to the model generator
 - And also to the migration generator

Checkpoint

What kind of object is likely being yielded in the migration code:

```
def up
  create_table :movies do |t|
    t.datetime 'release_date'
    # ...
  end
end
```

- (A) An object representing a database
- (B) An object representing an instance of a model
- (C) An object representing a table
- (D) Seriously, it could be anything

Checkpoint

What kind of object is likely being yielded in the migration code:

```
def up
  create_table :movies do |t|
    t.datetime 'release_date'
    # ...
  end
end
```

- (A) An object representing a database
- (B) An object representing an instance of a model
- (C) An object representing a table
- (D) Seriously, it could be anything

C — seriously, it's a table.

CRUD in SQL

- Structured Query Language (SQL) is the query language used by Relational Databases
- Rails **generates** SQL statements at runtime, based on Ruby code
- 4 basic operations on a table row: **C**reate, **R**ead, **U**ppdate, **D**eleate

The following is SQL (which we generally won't have to write)

```
INSERT INTO users(username, email, birthdate)
  VALUES ("joel", "jsommers@colgate.edu", "1972-06-10"),
  ("halima", "----", "2008-12-20")
SELECT * FROM users
  WHERE (birthdate BETWEEN "1990-01-01" AND "2010-01-01")
UPDATE users SET email = "jsommers@acm.org" WHERE username="joel"
DELETE FROM users WHERE id=1
```


Experimenting and debugging with a model

- It is often useful to “play” with a model class and its method
- `rails console` creates an irb session with your model classes preloaded
- Can do any CRUD operation on a model
- Awesome for experimenting...
- Also useful for debugging...
- See the SQL generated by any method...

The Ruby side of a model

- Models are subclassed from `ApplicationRecord`
 - "Connects" a model to the database
 - Provides CRUD operations on the model
- Rails convention: database table name is derived from model's name
 - `Product` (model class) \rightarrow `products` (table)
- Database table column names are getters and setters for model attributes
- **Key: the getters and setters do not simply modify instance variables!**

```
# 3 ways to create ActiveRecord objects
```

```
product = Product.new
product.name = 'Stuffed animal'
product.price = 11.99
```

```
product = Product.new do |p|
  p.name = 'Stuffed animal'
  p.price = 11.99
end
```

```
product = Product.new(:name => 'Stuffed animal', :price => 11.99)
```

Model attributes

- Model class provides a way to access database table row by manipulating Ruby object
- Table columns should generally be accessed using `self.attribute` within any instance methods of the model class
 - Why?

```
class Product < ApplicationRecord
  def capname
    self.name.capitalize
    # NOT @name.capitalize
  end
end
```

Creating a model object: `new` \neq `save`

- Must call `save` or `save!` on an ActiveRecord model instance to actually save changes to DB
 - `!` version is dangerous: throws an exception if operation fails
 - `create` method does both `new` and `save`
- Once created, object acquires a primary key (`id` column)
 - if `x.id` is `nil` or `x.new_record?` is true, `x` has never been saved
 - `x.persisted?` is another way to test whether an object has been saved to the database
 - These behaviors are inherited from `ActiveRecord::Base`

Read: finding things in the DB

`all` class method retrieves all objects

```
@products = Product.all
```

Other basic class methods for object retrieval:

- `first`: retrieves first object (lowest id)
- `last`: object with largest id

The `find` class method is used to find objects by id:

```
p = Product.find(13) # get object with id 13  
products = Product.find([13, 42]) # find understands arrays of ids
```

The `find_by` class method is used to find objects by attribute:

```
p = Product.find_by(name: "fuzzy bunny")
```

Read: where method

Class method `where` selects objects based on attributes

- Argument to `where` can be a string or a hash

```
Product.where('price > 2')
```

```
Product.where('price < ?', 2) # parameterized
```

```
Product.where('name like ?', "%#{matchstr}%")
```

```
# don't use Ruby string interpolation to fill in queries!
```

```
Product.where("price < #{limit}") # BAD! Do NOT do this!
```

`where` can be chained

```
Product.where('price <= 5').where('name like ?', "%#{fuzzy}%")
```

- http://guides.rubyonrails.org/active_record_basics.html

Checkpoint

Consider the following example, in which interpolation is performed in an ActiveRecord query:

```
Product.where("price > #{minprice}")
```

Why is interpolation bad in ActiveRecord queries?

- A. Interpolation doesn't work in this context; Ruby doesn't allow it.
- B. Performance; interpolation is too slow.
- C. It's a database security threat.
- D. It's bad because you should use `find_by` instead of `where`.

Checkpoint

Consider the following example, in which interpolation is performed in an ActiveRecord query:

```
Product.where("price > #{minprice}")
```

Why is interpolation bad in ActiveRecord queries?

- A. Interpolation doesn't work in this context; Ruby doesn't allow it.
 - B. Performance; interpolation is too slow.
 - C. It's a database security threat.
 - D. It's bad because you should use `find_by` instead of `where`.
- C. Where did `minprice` come from? If it came from your application, then maybe it's ok. If it came (even indirectly) from the browser, this is bad, bad, bad.

Updating

- Modify attributes, then save object

```
p = Product.find(42)
p.price = p.price * 1.1 # markup
p.save! # raises exception if save fails
p.save # returns true if saved, false otherwise
```

- Update attributes on existing object

```
Product.find(1).
  update(:price => 49.99) # accepts a hash
# or
Product.update(:id=>1, :price=>49.99)
```

- Update all matching records

```
# update_all takes hash of attributes to change or a string
Product.update_all("price = price * 0.90")
```

- Transactional: either all attributes are updated, or none are

Delete is straightforward

- Destroy is an instance method

```
p = Product.find(13)
p.destroy
```

- There's also `delete`, which does not trigger lifecycle callbacks (we'll discuss later)
 - Generally do **not** want `delete`; use `destroy`
- Once an AR object is destroyed, you can access but not modify the in-memory object

```
p.name = 'Slinky 2' # fails -
                    # object is 'frozen'
```

Checkpoint

Assume table `fortune_cookie` has column `fortune_text`. Assume also that we have a model `FortuneCookie < ApplicationRecord`. Which of these instance methods will not return a silly fortune (if any)?

- (A)

```
def silly_fortune_1
  @fortune_text + 'in bed'
end
```
- (B)

```
def silly_fortune_2
  self.fortune_text + 'in bed'
end
```
- (C)

```
def silly_fortune_3
  fortune_text + 'in bed'
end
```
- (D) They will all return a silly fortune

Checkpoint

Assume table `fortune_cookie` has column `fortune_text`. Assume also that we have a model `FortuneCookie < ApplicationRecord`. Which of these instance methods will not return a silly fortune (if any)?

(A)

```
def silly_fortune_1  
  @fortune_text + 'in bed'  
end
```

(B)

```
def silly_fortune_2  
  self.fortune_text + 'in bed'  
end
```

(C)

```
def silly_fortune_3  
  fortune_text + 'in bed'  
end
```

(D) They will all return a silly fortune

B or C are ok, but A is not. (Correct answer is A: it will not return a silly fortune.)

“Seeding” the database

- It's often useful to “seed” the database to provide an initial set of model instances (by adding rows to db tables)
- Add to `db/seeds.rb`
 - Use any ActiveRecord calls to create/save model instances
 - Load seeds into db with `rails db:seed`

```
# assuming decimal field for price
```

```
Product.create!(name: "fuzzy bunny", price: 11.99)
```

```
Product.create!(name: "slinky", price: 3.49)
```

```
# - or - make a hash of attributes, and use a each "loop"
```

```
products = [  
  { name: "fuzzy bunny", price: 11.99 },  
  { name: "slinky", price: 3.49 },  
]
```

```
products.each do |pdata|  
  Product.create!(pdata)  
end
```

Summary: ActiveRecord

- Subclassing `ApplicationRecord` (which derives directly from `ActiveRecord::Base`) “connects” a model to the database
- **C** (save, create), **R** (all, first, last, where, find), **U** (update, update_all), **D** (destroy)
- Convention over configuration maps
 - model name to DB table name
 - getters/setters to DB table columns
- Object in memory \neq row in database
 - `save` must be used to persist
 - `destroy` doesn't destroy in-memory copy